

## Assignment 2 - authored by, Revathy V

### 1. Download the dataset from the source here.

#### About the dataset:

This dataset is all about churn modelling of a credit company. It has the details about the end user who are using credit card and also it has some variables to depict the churn of the customer.

**RowNumber** - Serial number of the rows **CustomerId** - Unique identification of customer  
**Surname** - Name of the customer **CreditScore** - Cipil score of the customer **Geography** - Location of the bank **Gender** - Sex of the customer **Age** - Age of the customer **Tenure** - Repayment period for the credit amount **Balance** - Current balance in thier credit card  
**NumOfProducts** - Products owned by the customer from the company **HasCrCard** - Has credit card or not (0 - no , 1 - yes) **IsactiveMember** - Is a active member or not  
**EstimatedSalary** - Salary of the customer **Exited** - Churn of the customer

```
import warnings
warnings.filterwarnings("ignore")
```

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

### 2. Load the dataset

```
df = pd.read_csv("Churn_Modelling.csv")
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	
0	2	0.00	1	1	1	

1	1	83807.86	1	0	1
2	8	159660.80	3	1	0
3	1	0.00	2	0	0
4	2	125510.82	1	1	1

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

df.tail()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
9995	9996	15606229	Obijiaku	771	France	Male
39						
9996	9997	15569892	Johnstone	516	France	Male
35						
9997	9998	15584532	Liu	709	France	Female
36						
9998	9999	15682355	Sabbatini	772	Germany	Male
42						
9999	10000	15628319	Walker	792	France	Female
28						

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
9995	5	0.00	2	1		0
9996	10	57369.61	1	1		1
9997	7	0.00	1	0		1
9998	3	75075.31	2	1		0
9999	4	130142.79	1	1		0

	EstimatedSalary	Exited
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

### 3 a). Univariate analysis

*#checking for categorical variables*

```
category = df.select_dtypes(include=[np.object])
print("Categorical Variables: ",category.shape[1])
```

*#checking for numerical variables*

```
numerical = df.select_dtypes(include=[np.int64,np.float64])
print("Numerical Variables: ",numerical.shape[1])
```

```
Categorical Variables: 3  
Numerical Variables: 11
```

```
df.columns
```


```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore',  
      'Geography',  
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',  
      'HasCrCard',  
      'IsActiveMember', 'EstimatedSalary', 'Exited'],  
      dtype='object')
```

```
df.shape
```

```
(10000, 14)
```


```
credit = df['CreditScore']  
credit.plot(kind="hist",figsize=(8,5))
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff75105c850>
```




```
geo = df['Geography'].value_counts()  
geo.plot(kind="pie",figsize=(10,8))
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff750f68910>
```




```
sns.countplot(df['Gender'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff750f74090>
```




```
sns.distplot(df['Age'],hist=False)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff750a7b710>
```




```
plt.figure(figsize=(10,8))  
sns.countplot(df['Tenure'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff750978d90>
```




```
product = df['NumOfProducts'].value_counts()  
product.plot(kind="pie",figsize=(10,8))
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7ff750908390>




```
cr = df['HasCrCard'].value_counts()
cr.plot(kind="pie",figsize=(10,8))
<matplotlib.axes._subplots.AxesSubplot at 0x7ff7508e2fd0>
```




```
plt.figure(figsize=(8,5))
sns.countplot(df['IsActiveMember'])
<matplotlib.axes._subplots.AxesSubplot at 0x7ff75089dad0>
```






```
sns.distplot(df['EstimatedSalary'],hist=False)  
<matplotlib.axes._subplots.AxesSubplot at 0x7ff750806c10>
```



```
plt.figure(figsize=(8,5))  
sns.countplot(df['Exited'])
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7ff750781bd0>




### Inference:

1. The data has 11 numerical variables and 3 categorical variables.
2. It has 10000 rows and 14 columns
3. The normalized credit score is around 700, More than 500 people have credit score greater than 800.
4. France occupies 50% of customers, where as Germany and Spain shared equal.
5. Dataset is dominated by Male Customers.
6. Median age is around 40 to 45.
7. Highest number of customer has thier tenure period for 2 years.
8. Credit company has maximum customers, who uses single product.
9. Most of the customer has credit card.
10. More than 40% of the population is not an active member.
11. The Churn is less compared to the satisfaction. **Dataset is imbalanced.**


### 3 b). Bivariate analysis

```
sns.barplot(x='Gender',y='CreditScore',hue='Geography',data=df)
```


<matplotlib.axes.\_subplots.AxesSubplot at 0x7ff75075b190>




```
sns.violinplot(x='Geography',y='Balance',data=df)  
<matplotlib.axes._subplots.AxesSubplot at 0x7ff75066cbd0>
```




```
sns.violinplot(x='Geography',y='EstimatedSalary',data=df)  
<matplotlib.axes._subplots.AxesSubplot at 0x7ff7506016d0>
```




```
sns.violinplot(x='Gender',y='Balance',data=df)  
<matplotlib.axes._subplots.AxesSubplot at 0x7ff75057f410>
```




```
sns.barplot(x='Exited',y='CreditScore',hue='Gender',data=df)  
<matplotlib.axes._subplots.AxesSubplot at 0x7ff7519dd810>
```



```
sns.barplot(x='Exited',y='CreditScore',hue='Geography',data=df)
<matplotlib.axes._subplots.AxesSubplot at 0x7ff7504dc2d0>
```




```
sns.barplot(x='IsActiveMember',y='EstimatedSalary',hue='Gender',data=d
f)
<matplotlib.axes._subplots.AxesSubplot at 0x7ff750983590>
```



```
sns.barplot(x='Exited',y='Tenure',data=df)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7ff750680750>
```



### Inference:

1. Credit score for Male is higher in Spain.


2. Average bank salary lies in the range of 100k to 150k.
3. Estimated salary is normalized and same for all country.
4. Credit score for churn is low.
5. Churn in Germany is higher compared to other countries.
6. Exited people tenure period is around 6 years.

### 3 c). Multivariate analysis

```
gp1 = df.groupby('Gender')['Geography'].value_counts()  
gp1.plot(kind='pie',figsize=(10,8))  
print(gp1)
```

Gender	Geography	
Female	France	2261
	Germany	1193
	Spain	1089
Male	France	2753
	Spain	1388
	Germany	1316

Name: Geography, dtype: int64



```
gp2 = df.groupby('Gender')['Age'].mean()
print(gp2)
```

```
Gender
Female    39.238389
Male      38.658237
Name: Age, dtype: float64
```

```
gp3 = df.groupby(['Gender', 'Geography'])['Tenure'].mean()
print(gp3)
```


```
Gender  Geography
Female  France      4.950022
        Germany     4.965633
        Spain      5.000000
Male    France      5.049401
        Germany     5.050152
        Spain      5.057637
Name: Tenure, dtype: float64
```



```
gp4 = df.groupby(['HasCrCard', 'IsActiveMember'])
['Geography'].value_counts()
gp4.plot(kind='bar', figsize=(8,5))
print(gp4)
```

HasCrCard	IsActiveMember	Geography	
0	0	France	706
		Spain	352
		Germany	343
	1	France	765
		Spain	404
		Germany	375
1	0	France	1717
		Germany	918
		Spain	813
	1	France	1826
		Spain	908
		Germany	873


Name: Geography, dtype: int64



```
gp5 = df.groupby(['Gender', 'HasCrCard', 'IsActiveMember'])
['EstimatedSalary'].mean()
gp5.plot(kind="line", figsize=(10,8))
print(gp5)
```

Gender	HasCrCard	IsActiveMember	
Female	0	0	102006.080352
		1	102648.996944
	1	0	101208.014567
		1	98510.152300
Male	0	0	99756.431151
		1	99873.931251
	1	0	100353.378996
		1	98914.378703


Name: EstimatedSalary, dtype: float64



```
gp6 = df.groupby(['Gender', 'IsActiveMember'])['Exited'].value_counts()
gp6.plot(kind='bar', figsize=(10,8))
print(gp6)
```

Gender	IsActiveMember	Exited	
Female	0	0	1534
		1	725
	1	0	1870
		1	414
Male	0	0	2013
		1	577
	1	0	2546
		1	

1 321  
Name: Exited, dtype: int64



```
gp7 = df.groupby('Exited')['Balance', 'EstimatedSalary'].mean()
print(gp7)
```


	Balance	EstimatedSalary
Exited		
0	72745.296779	99738.391772
1	91108.539337	101465.677531

```
gp8 = df.groupby(['Geography', 'Exited'])['Gender'].value_counts()
gp8.plot(kind='bar', figsize=(10,8))
print (gp8)
```

Geography	Exited	Gender	
France	0	Male	2403
		Female	1801
	1	Female	460
		Male	350
Germany	0	Male	950

		Female	745
	1	Female	448
		Male	366
Spain	0	Male	1206
		Female	858
	1	Female	231
		Male	182

Name: Gender, dtype: int64



## Inference:

1. Germany has more female customers compared to male customers.
2. Average age of Male is 38, whereas average age of Female is 39.
3. Tenure period for both male and female is high in Spain.
4. It is observed that, those who have credit card are very active member in the company.

5. The estimated salary for a person who is not having credit card is high when compared to those having them.
6. Churn for inactive member is high compared to active member.
7. Those who churn has thier estimated salary very low.
8. France has the more churn rate.

## 4. Descriptive statistics

df.describe().T

	count	mean	std	min \
RowNumber	10000.0	5.000500e+03	2886.895680	1.00
CustomerId	10000.0	1.569094e+07	71936.186123	15565701.00
CreditScore	10000.0	6.505288e+02	96.653299	350.00
Age	10000.0	3.892180e+01	10.487806	18.00
Tenure	10000.0	5.012800e+00	2.892174	0.00
Balance	10000.0	7.648589e+04	62397.405202	0.00
NumOfProducts	10000.0	1.530200e+00	0.581654	1.00
HasCrCard	10000.0	7.055000e-01	0.455840	0.00
IsActiveMember	10000.0	5.151000e-01	0.499797	0.00
EstimatedSalary	10000.0	1.000902e+05	57510.492818	11.58
Exited	10000.0	2.037000e-01	0.402769	0.00

	25%	50%	75%	max
RowNumber	2500.75	5.000500e+03	7.500250e+03	10000.00
CustomerId	15628528.25	1.569074e+07	1.575323e+07	15815690.00
CreditScore	584.00	6.520000e+02	7.180000e+02	850.00
Age	32.00	3.700000e+01	4.400000e+01	92.00
Tenure	3.00	5.000000e+00	7.000000e+00	10.00
Balance	0.00	9.719854e+04	1.276442e+05	250898.09
NumOfProducts	1.00	1.000000e+00	2.000000e+00	4.00
HasCrCard	0.00	1.000000e+00	1.000000e+00	1.00
IsActiveMember	0.00	1.000000e+00	1.000000e+00	1.00
EstimatedSalary	51002.11	1.001939e+05	1.493882e+05	199992.48
Exited	0.00	0.000000e+00	0.000000e+00	1.00

## 5. Handling the missing values

```
df.isnull().sum()
```



```
RowNumber      0
CustomerId     0
Surname        0
CreditScore    0
Geography      0
Gender         0
Age           0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard      0
IsActiveMember 0
EstimatedSalary 0
Exited         0
dtype: int64
```

**There is no missing value in the dataset**

## 6. Finding outliers



```
def replace_outliers(df, field_name):
    Q1 = np.percentile(df[field_name],25,interpolation='midpoint')
    Q3 = np.percentile(df[field_name],75,interpolation='midpoint')
    IQR = Q3-Q1
    maxi = Q3+1.5*IQR
    mini = Q1-1.5*IQR
    df[field_name]=df[field_name].mask(df[field_name]>maxi,maxi)
    df[field_name]=df[field_name].mask(df[field_name]<mini,mini)

plt.title("Before removing outliers")
sns.boxplot(df['CreditScore'])
plt.show()
plt.title("After removing outliers")
replace_outliers(df, 'CreditScore')
sns.boxplot(df['CreditScore'])
plt.show()
```



```
plt.title("Before removing outliers")
sns.boxplot(df['Age'])
plt.show()
plt.title("After removing outliers")
replace_outliers(df, 'Age')
```


```
sns.boxplot(df['Age'])  
plt.show()
```



```
sns.boxplot(df['Balance'])
```




<matplotlib.axes.\_subplots.AxesSubplot at 0x7ff74ff93a50>



```
sns.boxplot(df['EstimatedSalary'])
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7ff74fe1ae50>



**Outliers from Age and Credit Score columns are removed**

## 7. Check for categorical column and perform encoding.

```
from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()
```

```
df['Gender'] = le.fit_transform(df['Gender'])  
df['Geography'] = le.fit_transform(df['Geography'])
```

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
0	1	15634602	Hargrave	619.0	0	0
42.0						
1	2	15647311	Hill	608.0	2	0
41.0						
2	3	15619304	Onio	502.0	0	0
42.0						
3	4	15701354	Boni	699.0	0	0
39.0						
4	5	15737888	Mitchell	850.0	2	0
43.0						

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

**Only two columns(Gender and Geography) is label encoded**

## Removing unwanted columns and checking for feature importance

```
df = df.drop(['RowNumber', 'CustomerId', 'Surname'],axis=1)
```

```
df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance
NumOfProducts \						
0	619.0	0	0	42.0	2	0.00
1						
1	608.0	2	0	41.0	1	83807.86
1						

```

2          502.0          0          0  42.0          8  159660.80
3
3          699.0          0          0  39.0          1          0.00
2
4          850.0          2          0  43.0          2  125510.82
1

```

```

      HasCrCard  IsActiveMember  EstimatedSalary  Exited
0             1             1         101348.88         1
1             0             1         112542.58         0
2             1             0         113931.57         1
3             0             0          93826.63         0
4             1             1          79084.10         0


```

```

plt.figure(figsize=(20,10))
df_lt = df.corr(method = "pearson")
df_lt1 = df_lt.where(np.tril(np.ones(df_lt.shape)).astype(np.bool))
sns.heatmap(df_lt1,annot=True,cmap="coolwarm")

```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7ff74f5ba590>



**1. The Removed columns are nothing to do with model building. 2. Feature importance also checked using pearson correlation.**

## 8. Data Splitting

```

target = df['Exited']
data = df.drop(['Exited'],axis=1)

print(data.shape)
print(target.shape)

```

```
(10000, 10)
(10000,)
```

## 9. Scaling the independent values

```
from sklearn.preprocessing import StandardScaler
se = StandardScaler()

data['CreditScore'] =
se.fit_transform(pd.DataFrame(data['CreditScore']))
data['Age'] = se.fit_transform(pd.DataFrame(data['Age']))
data['Balance'] = se.fit_transform(pd.DataFrame(data['Balance']))
data['EstimatedSalary'] =
se.fit_transform(pd.DataFrame(data['EstimatedSalary']))

data.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance
0	-0.326878	0	0	0.342615	2	-1.225848
1	-0.440804	2	0	0.240011	1	0.117350
2	-1.538636	0	0	0.342615	8	1.333053
3	0.501675	0	0	0.034803	1	-1.225848
4	2.065569	2	0	0.445219	2	0.785728

	HasCrCard	IsActiveMember	EstimatedSalary
0	1	1	0.021886
1	0	1	0.216534
2	1	0	0.240687
3	0	0	-0.108918
4	1	1	-0.365276

## 10. Train test split

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =
train_test_split(data,target,test_size=0.25,random_state=101)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(7500, 10)
(2500, 10)
```

(7500,)  
(2500,)

### Conclusion:

1. The model is scaled using StandardScaler method.
2. The train and test split ratio is 15:5.
3. As it is a classification problem, basic algorithms can be used to build ML models.